

# CIND820\_Capstone\_Project

June 26, 2023

## 0.1 CIND820 - Capstone Project

### 1 Investigate Airline passenger satisfaction using Machine Learning Techniques

### 2 Preparation:

```
[ ]: ! python -V
```

Python 3.10.12

Import csv file (the dataset and the data dictionary)

```
[1]: # Importing required libraries

import pandas as pd

import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn import preprocessing
```

```
[2]: # Import the dataset
# Use first column "ID" as Index by using index_col=0

url = 'https://raw.githubusercontent.com/HitomiMo/CIND820_Capstone-Project/main/
↳airline_passenger_satisfaction.csv'
df1 = pd.read_csv(url, index_col=0)
df1.head()
```

```
[2]:
```

	Gender	Age	Customer Type	Type of Travel	Class	Flight Distance	\
ID							
1	Male	48	First-time	Business	Business	821	
2	Female	35	Returning	Business	Business	821	
3	Male	41	Returning	Business	Business	853	

4	Male	50	Returning	Business	Business	1905
5	Female	49	Returning	Business	Business	3470

	Departure Delay	Arrival Delay	Departure and Arrival Time Convenience
ID			
1	2	5.0	3
2	26	39.0	2
3	0	0.0	4
4	0	0.0	2
5	0	1.0	3

	Ease of Online Booking	On-board Service	Seat Comfort
ID			
1	3	3	5
2	2	5	4
3	4	3	5
4	2	5	5
5	3	3	4

	Leg Room Service	Cleanliness	Food and Drink	In-flight Service
ID				
1	2	5	5	5
2	5	5	3	5
3	3	5	5	3
4	5	4	4	5
5	4	5	4	3

	In-flight Wifi Service	In-flight Entertainment	Baggage Handling
ID			
1	3	5	5
2	2	5	5
3	4	3	3
4	2	5	5
5	3	3	3

	Satisfaction
ID	
1	Neutral or Dissatisfied
2	Satisfied
3	Satisfied
4	Satisfied
5	Satisfied

[5 rows x 23 columns]

```
[ ]: # Import the dictionary
```

```
url2 = 'https://raw.githubusercontent.com/HitomiMo/CIND820_Capstone-Project/
↪main/data_dictionary.csv'
data_dictionary = pd.read_csv(url2, index_col=0)
data_dictionary
```

```
[ ]: Description
Field
ID Unique passenger
identifier
Gender Gender of the passenger
(Female/Male)
Age Age of the
passenger
Customer Type Type of airline customer (First-
time/Returning)
Type of Travel Purpose of the flight
(Business/Personal)
Class Travel class in the airplane for the
passenger...
Flight Distance Flight distance
in miles
Departure Delay Flight departure delay
in minutes
Arrival Delay Flight arrival delay
in minutes
Departure and Arrival Time Convenience Satisfaction level with the convenience
of the...
Ease of Online Booking Satisfaction level with the online
booking exp...
Check-in Service Satisfaction level with the check-in
service f...
Online Boarding Satisfaction level with the online
boarding ex...
Gate Location Satisfaction level with the gate
location in t...
On-board Service Satisfaction level with the on-boarding
servic...
Seat Comfort Satisfaction level with the comfort of
the air...
Leg Room Service Satisfaction level with the leg room of
the ai...
Cleanliness Satisfaction level with the cleanliness
of the...
Food and Drink Satisfaction level with the food and
drinks on...
In-flight Service Satisfaction level with the in-flight
service ...
```

In-flight Wifi Service	Satisfaction level with the in-flight
Wifi ser...	
In-flight Entertainment	Satisfaction level with the in-flight
entertai...	
Baggage Handling	Satisfaction level with the baggage
handling f...	
Satisfaction	Overall satisfaction level with the
airline (S...	

### 3 Exploratory Data Analysis (EDA)

Install pandas-profiling

```
[ ]: pip install pandas-profiling
```

```
[ ]: from pandas_profiling import ProfileReport
prof = ProfileReport(df1)
prof.to_file(output_file='output.html')
```

Summarize dataset: 0%| | 0/5 [00:00<?, ?it/s]

Generate report structure: 0%| | 0/1 [00:00<?, ?it/s]

Render HTML: 0%| | 0/1 [00:00<?, ?it/s]

Export report to file: 0%| | 0/1 [00:00<?, ?it/s]

```
[ ]: df1.head(10)
```

```
[ ]:
Gender  Age Customer Type Type of Travel      Class  Flight Distance \
ID
1    Male   48   First-time      Business  Business           821
2  Female   35   Returning      Business  Business           821
3    Male   41   Returning      Business  Business           853
4    Male   50   Returning      Business  Business          1905
5  Female   49   Returning      Business  Business          3470
6    Male   43   Returning      Business  Business          3788
7    Male   43   Returning      Business  Business          1963
8  Female   60   Returning      Business  Business           853
9    Male   50   Returning      Business  Business          2607
10 Female   38   Returning      Business  Business          2822

Departure Delay  Arrival Delay  Departure and Arrival Time Convenience \
ID
1                2             5.0                      3
2               26            39.0                      2
3                0             0.0                      4
4                0             0.0                      2
```

5	0	1.0	3
6	0	0.0	4
7	0	0.0	3
8	0	3.0	3
9	0	0.0	1
10	13	0.0	2

	Ease of Online Booking	...	On-board Service	Seat Comfort	\
ID		...			
1	3	...	3	5	
2	2	...	5	4	
3	4	...	3	5	
4	2	...	5	5	
5	3	...	3	4	
6	4	...	4	4	
7	3	...	5	5	
8	4	...	3	4	
9	1	...	4	3	
10	5	...	5	4	

	Leg Room Service	Cleanliness	Food and Drink	In-flight Service	\
ID					
1	2	5	5	5	
2	5	5	3	5	
3	3	5	5	3	
4	5	4	4	5	
5	4	5	4	3	
6	4	3	3	4	
7	5	4	5	5	
8	4	4	4	3	
9	4	3	3	4	
10	5	4	2	5	

	In-flight Wifi Service	In-flight Entertainment	Baggage Handling	\
ID				
1	3	5	5	
2	2	5	5	
3	4	3	3	
4	2	5	5	
5	3	3	3	
6	4	4	4	
7	3	5	5	
8	4	3	3	
9	4	4	4	
10	2	5	5	

Satisfaction

```
ID
1  Neutral or Dissatisfied
2              Satisfied
3              Satisfied
4              Satisfied
5              Satisfied
6              Satisfied
7              Satisfied
8              Satisfied
9  Neutral or Dissatisfied
10             Satisfied
```

[10 rows x 23 columns]

Check the dataset

```
[ ]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 129880 entries, 1 to 129880
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                129880 non-null  object
1   Age                                  129880 non-null  int64
2   Customer Type                        129880 non-null  object
3   Type of Travel                       129880 non-null  object
4   Class                                129880 non-null  object
5   Flight Distance                      129880 non-null  int64
6   Departure Delay                     129880 non-null  int64
7   Arrival Delay                       129487 non-null  float64
8   Departure and Arrival Time Convenience 129880 non-null  int64
9   Ease of Online Booking               129880 non-null  int64
10  Check-in Service                     129880 non-null  int64
11  Online Boarding                      129880 non-null  int64
12  Gate Location                       129880 non-null  int64
13  On-board Service                     129880 non-null  int64
14  Seat Comfort                         129880 non-null  int64
15  Leg Room Service                     129880 non-null  int64
16  Cleanliness                          129880 non-null  int64
17  Food and Drink                       129880 non-null  int64
18  In-flight Service                    129880 non-null  int64
19  In-flight Wifi Service               129880 non-null  int64
20  In-flight Entertainment              129880 non-null  int64
21  Baggage Handling                     129880 non-null  int64
22  Satisfaction                         129880 non-null  object
dtypes: float64(1), int64(17), object(5)
memory usage: 23.8+ MB
```

**Observation:** \* Number of variable: 23 \* Number of entries: 129880

Check missing data

```
[ ]: missing_values = pd.isnull(df1)
missing_values.head()
```

```
[ ]:
Gender    Age    Customer Type    Type of Travel    Class    Flight Distance \
ID
1    False    False                False                False    False                False
2    False    False                False                False    False                False
3    False    False                False                False    False                False
4    False    False                False                False    False                False
5    False    False                False                False    False                False
```

```
Departure Delay    Arrival Delay    Departure and Arrival Time Convenience \
ID
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False
5                False                False                False
```

```
Ease of Online Booking    ...    On-board Service    Seat Comfort \
ID    ...
1                False    ...                False                False
2                False    ...                False                False
3                False    ...                False                False
4                False    ...                False                False
5                False    ...                False                False
```

```
Leg Room Service    Cleanliness    Food and Drink    In-flight Service \
ID
1                False                False                False                False
2                False                False                False                False
3                False                False                False                False
4                False                False                False                False
5                False                False                False                False
```

```
In-flight Wifi Service    In-flight Entertainment    Baggage Handling \
ID
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False
5                False                False                False
```

Satisfaction

```
ID
1      False
2      False
3      False
4      False
5      False
```

[5 rows x 23 columns]

```
[ ]: df1.isnull().sum()
```

```
[ ]: Gender          0
Age                0
Customer Type      0
Type of Travel     0
Class              0
Flight Distance    0
Departure Delay    0
Arrival Delay      393
Departure and Arrival Time Convenience  0
Ease of Online Booking  0
Check-in Service   0
Online Boarding    0
Gate Location      0
On-board Service   0
Seat Comfort       0
Leg Room Service   0
Cleanliness        0
Food and Drink     0
In-flight Service  0
In-flight Wifi Service  0
In-flight Entertainment  0
Baggage Handling   0
Satisfaction       0
dtype: int64
```

Check description of the data

```
[ ]: df1.describe()
```

```
[ ]:
count    129880.000000    129880.000000    129880.000000    129487.000000 \
mean       39.427957      1190.316392        14.713713        15.091129
std        15.119360       997.452477        38.071126        38.465650
min         7.000000        31.000000         0.000000         0.000000
25%        27.000000       414.000000         0.000000         0.000000
50%        40.000000       844.000000         0.000000         0.000000
```



75%	51.000000	1744.000000	12.000000	13.000000
max	85.000000	4983.000000	1592.000000	1584.000000

	Departure and Arrival Time Convenience	Ease of Online Booking \
count	129880.000000	129880.000000
mean	3.057599	2.756876
std	1.526741	1.401740
min	0.000000	0.000000
25%	2.000000	2.000000
50%	3.000000	3.000000
75%	4.000000	4.000000
max	5.000000	5.000000

	Check-in Service	Online Boarding	Gate Location	On-board Service \
count	129880.000000	129880.000000	129880.000000	129880.000000
mean	3.306267	3.252633	2.976925	3.383023
std	1.266185	1.350719	1.278520	1.287099
min	0.000000	0.000000	0.000000	0.000000
25%	3.000000	2.000000	2.000000	2.000000
50%	3.000000	3.000000	3.000000	4.000000
75%	4.000000	4.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000	5.000000

	Seat Comfort	Leg Room Service	Cleanliness	Food and Drink \
count	129880.000000	129880.000000	129880.000000	129880.000000
mean	3.441361	3.350878	3.286326	3.204774
std	1.319289	1.316252	1.313682	1.329933
min	0.000000	0.000000	0.000000	0.000000
25%	2.000000	2.000000	2.000000	2.000000
50%	4.000000	4.000000	3.000000	3.000000
75%	5.000000	4.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000	5.000000

	In-flight Service	In-flight Wifi Service	In-flight Entertainment \
count	129880.000000	129880.000000	129880.000000
mean	3.642193	2.728696	3.358077
std	1.176669	1.329340	1.334049
min	0.000000	0.000000	0.000000
25%	3.000000	2.000000	2.000000
50%	4.000000	3.000000	4.000000
75%	5.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000

	Baggage Handling
count	129880.000000
mean	3.632114
std	1.180025

```

min            1.000000
25%            3.000000
50%            4.000000
75%            5.000000
max            5.000000

```

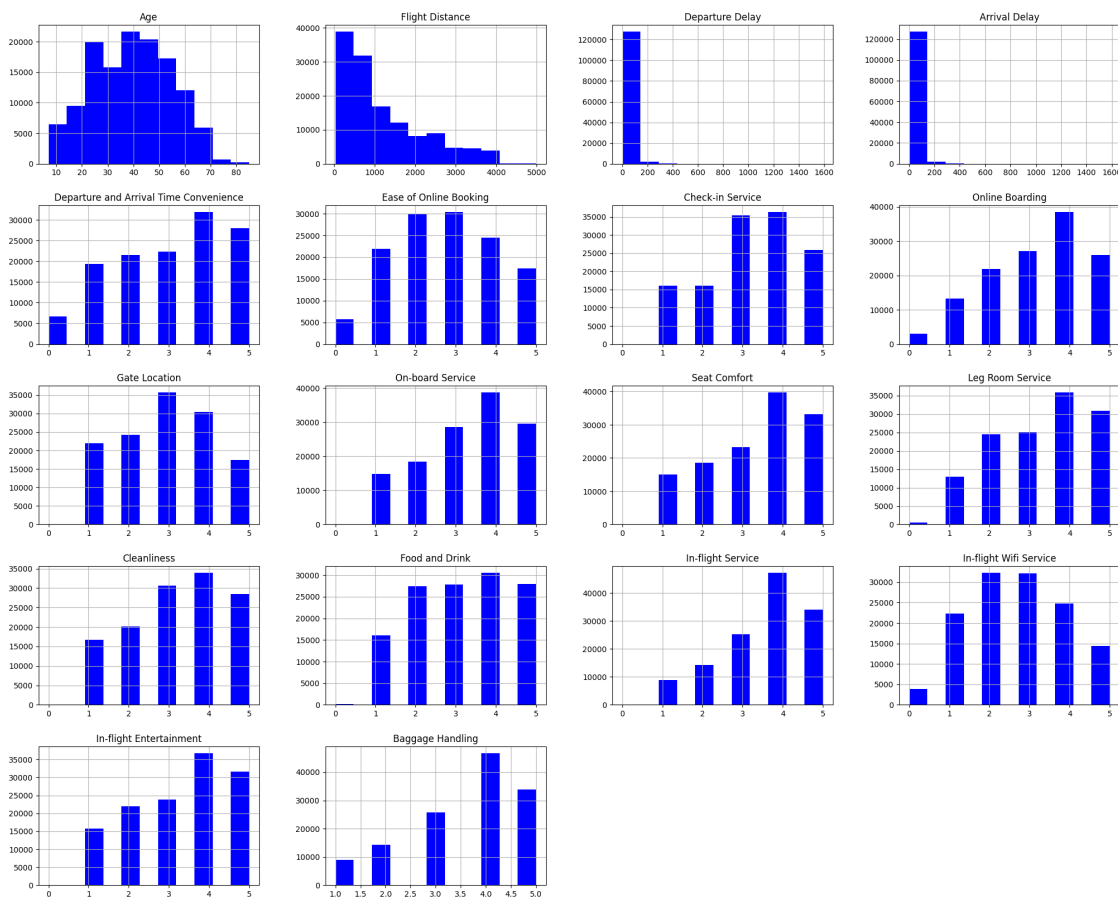
Check distribution of numerical variables in histogram

```

[ ]: # import numpy as np
     # import seaborn as sns
     # from matplotlib import pyplot as plt

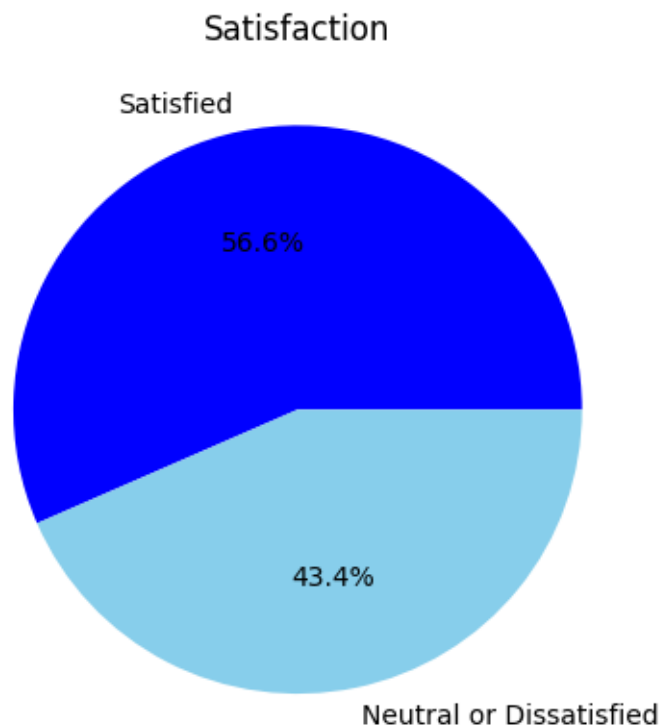
[ ]: binwidth=2
     df1.iloc[:,1:].hist(bins=11, figsize=(25,20), color='blue')
     plt.show()

```




Check categorical variables

```
[ ]: y = df1["Satisfaction"].value_counts()
labels = ["Satisfied", "Neutral or Dissatisfied"]
mycolors = ["blue", "skyblue"]
y = df1["Satisfaction"].value_counts()
plt.pie(y, labels = labels, colors = mycolors, autopct = '%1.1f%%')
plt.title('Satisfaction')
plt.show()
```



```
[ ]: # Check the # of customers - Satisfaction
mycolors2 = ["blue", "skyblue"]
s = sns.countplot(x='Satisfaction', data=df1)
abs_values = df1['Satisfaction'].value_counts().values

s.bar_label(container=s.containers[0], labels=abs_values);
```




**Observation:**

- Satisfaction is a target class and it is imbalanced.


Pie chart for Gender

```
[ ]: y = df1["Gender"].value_counts()
      labels = ["Male", "Female"]
      mycolors = ["Green", "Yellow"]
      plt.pie(y, labels = labels, colors=mycolors, autopct='%1.1f%%')
      plt.show()
```



```
[ ]: # Check the # of customers - Gender
s = sns.countplot(x='Gender',data=df1)
abs_values = df1['Satisfaction'].value_counts().values


s.bar_label(container=s.containers[0], labels=abs_values);
```



```
[ ]: # Grouping the data points based on Gender
df1.groupby('Gender')['Satisfaction'].value_counts()
```


```
[ ]: Gender  Satisfaction
      Female  Neutral or Dissatisfied    37630
      Female  Satisfied                28269
      Male   Neutral or Dissatisfied    35822
      Male   Satisfied                28159
      Name: Satisfaction, dtype: int64
```

```
[ ]: sns.countplot(data = df1, x= df1['Satisfaction'], hue = df1["Gender"]);
```



Pie chart for Customer Type

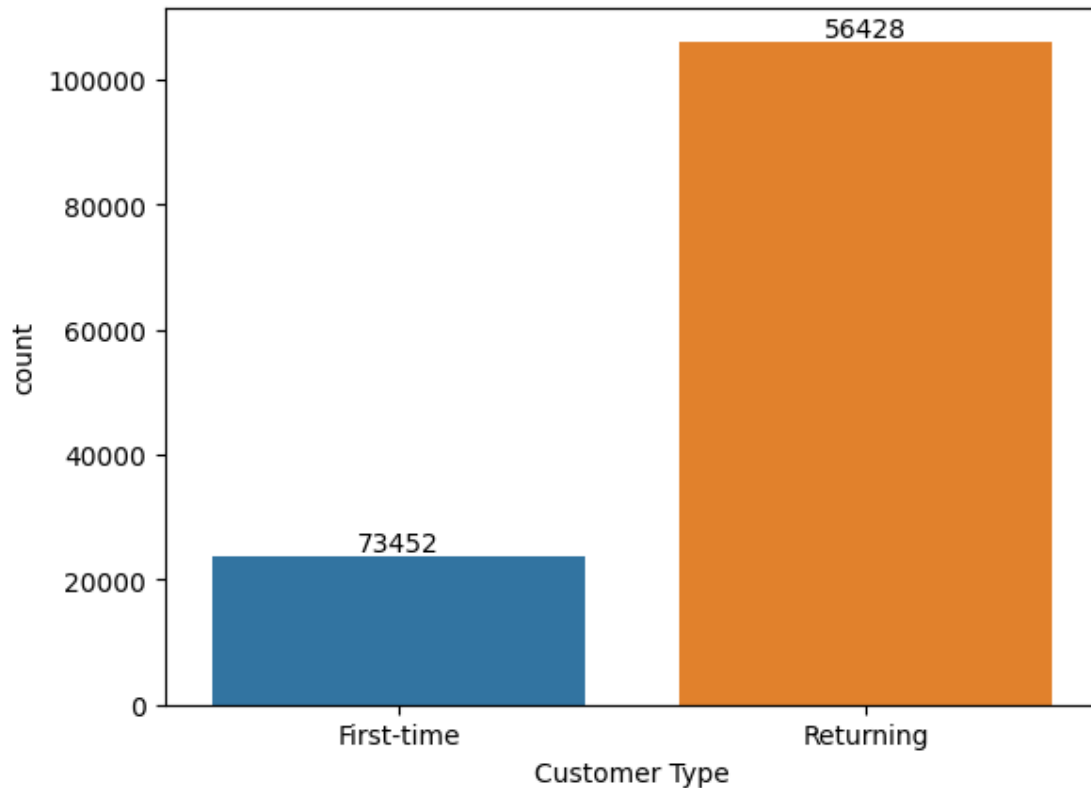
```
[ ]: y = df1["Customer Type"].value_counts()
labels = ["First-time", "Returning"]
mycolors = ["Brown", "Gray"]
plt.pie(y, labels = labels, colors=mycolors, autopct='%1.1f%%')
plt.show()
```



```
[ ]: # Check the # of customers - Customer Type
s = sns.countplot(x='Customer Type',data=df1)
abs_values = df1['Satisfaction'].value_counts().values

s.bar_label(container=s.containers[0], labels=abs_values);
```






```
[ ]: # Grouping the data points based on Customer Type
df1.groupby('Customer Type')['Satisfaction'].value_counts()
```


```
[ ]: Customer Type  Satisfaction
First-time      Neutral or Dissatisfied  18080
                Satisfied                5700
Returning       Neutral or Dissatisfied  55372
                Satisfied                50728
Name: Satisfaction, dtype: int64
```

```
[ ]: sns.countplot(data = df1, x= df1['Satisfaction'], hue = df1["Customer Type"]);
```




Pie chart for Type of Travel

```
[ ]: y = df1["Type of Travel"].value_counts()
labels = ["Business", "Personal"]
mycolors = ["Lightblue", "Lightgreen"]
plt.pie(y, labels = labels, colors=mycolors, autopct='%1.1f%%')
plt.show()
```



```
[ ]: # Check the # of customers - Customer Type
s = sns.countplot(x='Type of Travel',data=df1)
abs_values = df1['Satisfaction'].value_counts().values


s.bar_label(container=s.containers[0], labels=abs_values);
```



```
[ ]: # Grouping the data points based on Type of Travel
df1.groupby('Type of Travel')['Satisfaction'].value_counts()
```


```
[ ]: Type of Travel  Satisfaction
Business          Satisfied          52356
                  Neutral or Dissatisfied 37337
Personal          Neutral or Dissatisfied 36115
                  Satisfied           4072
Name: Satisfaction, dtype: int64
```

```
[ ]: sns.countplot(data = df1, x= df1['Satisfaction'], hue = df1["Type of Travel"]);
```



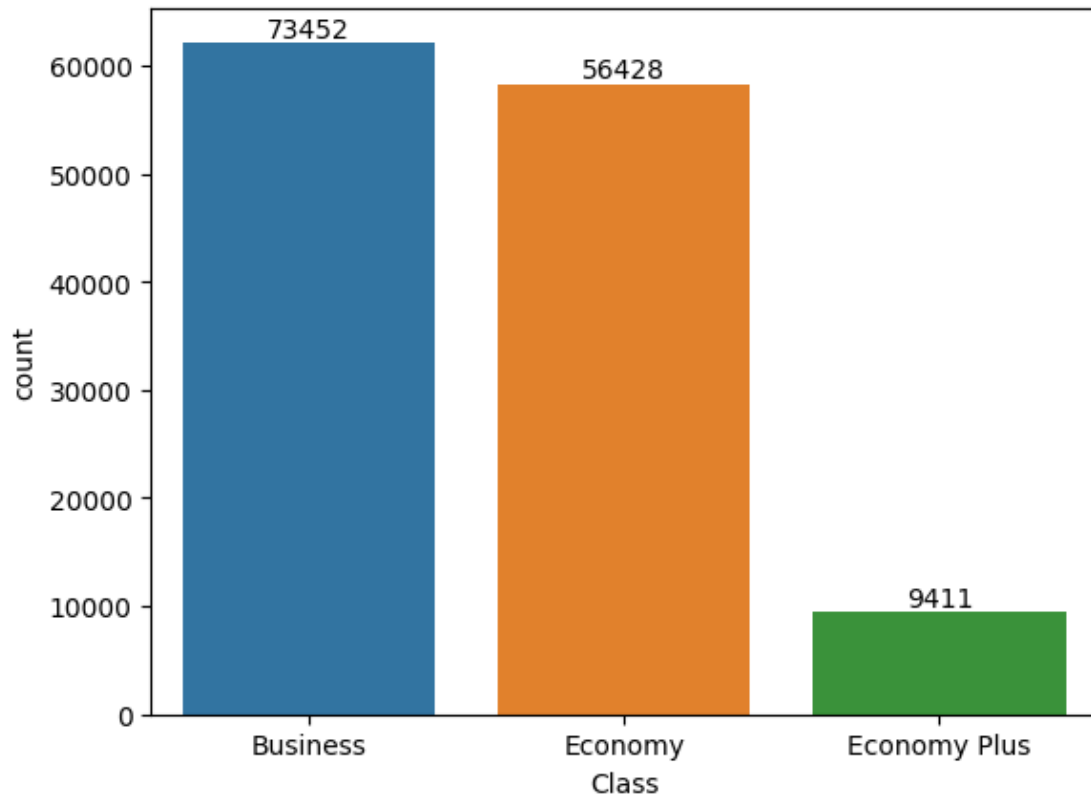
Pie chart for Class

```
[ ]: y = df1["Class"].value_counts()
      labels = ["Business", "Economy", "Economy Plus"]
      mycolors = ["Lightblue", "Lightgreen", "Lightyellow"]
      plt.pie(y, labels = labels, colors=mycolors, autopct='%1.1f%%')
      plt.show()
```



```
[ ]: # Check the # of customers - Customer Type
s = sns.countplot(x='Class',data=df1)
abs_values = df1['Satisfaction'].value_counts().values

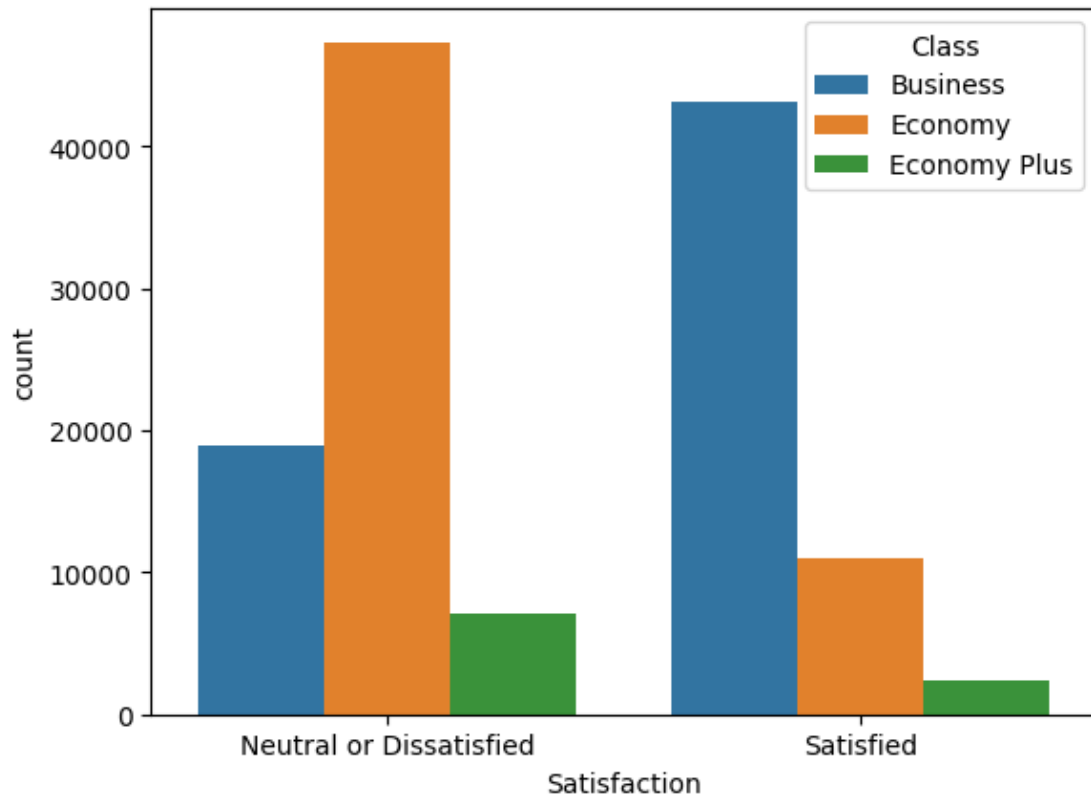
s.bar_label(container=s.containers[0], labels=abs_values);
```



```
[ ]: # Grouping the data points based on Class
df1.groupby('Class')['Satisfaction'].value_counts()
```

```
[ ]: Class      Satisfaction
      Business      Satisfied      43166
      Business  Neutral or Dissatisfied  18994
      Economy    Neutral or Dissatisfied  47366
      Economy      Satisfied      10943
      Economy Plus Neutral or Dissatisfied  7092
      Economy Plus      Satisfied      2319
      Name: Satisfaction, dtype: int64
```

```
[ ]: sns.countplot(data = df1, x= df1['Satisfaction'], hue = df1["Class"]);
```



Check correlation matrix


```
[ ]: plt.figure(figsize = (15,8))
      sns.heatmap(df1.corr() , annot = True , cmap = "YlGnBu")
```

<ipython-input-29-46ea086570be>:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
      sns.heatmap(df1.corr() , annot = True , cmap = "YlGnBu")
```

```
[ ]: <Axes: >
```





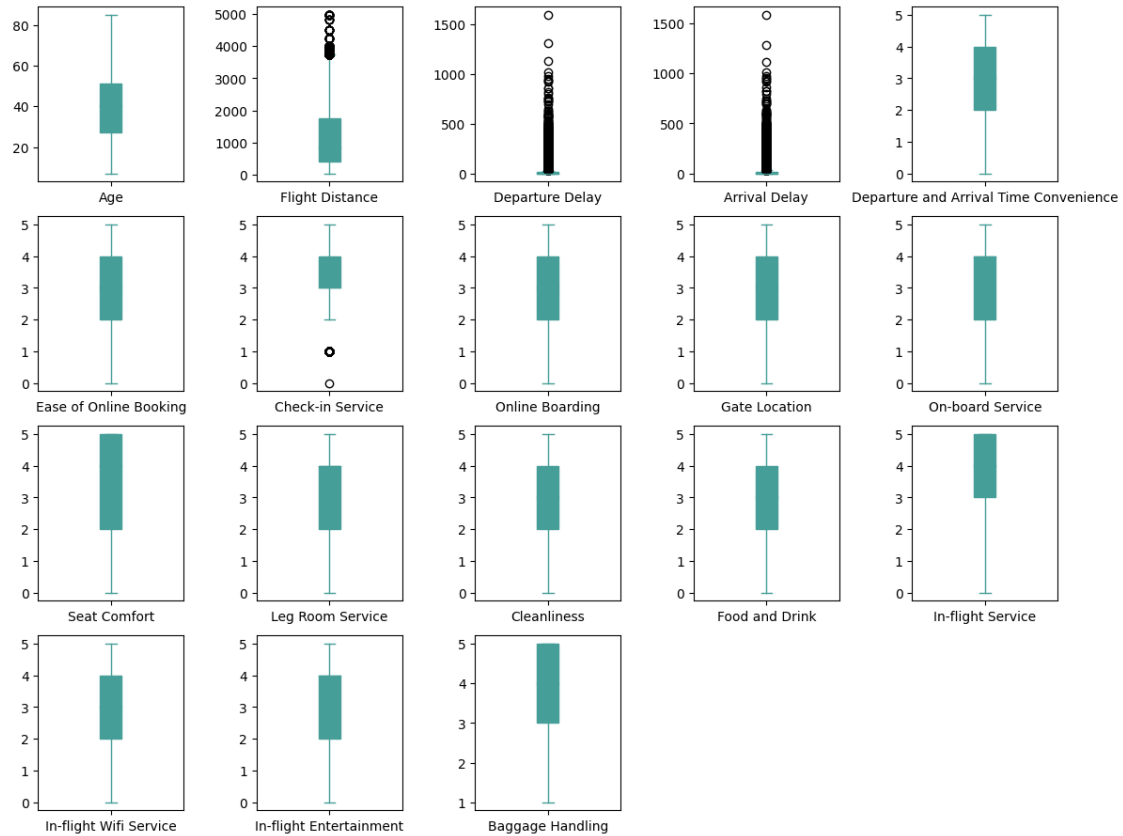
### Observation:

- **Strong positive correlation** between Arrival Delay and Departure Delay which is 0.97.
- **Moderate positive correlation** between In-flight Wi-Fi Service and Ease of Online Booking which is 0.71.


Check outliers of numerical variables in histogram

```
[ ]: # df1.plot(kind='box', subplots=True, layout=(8,5), figsize=(17,20))
```


```
[ ]: df1.plot(kind='box', subplots=True, layout=(5,5),
    ↪figsize=(14,14),patch_artist=True,color="#459E97")
plt.subplots_adjust(wspace = 0.5)
```




```
[ ]: # check distribution of Flight Distance
plt.figure(figsize = (5,5))
sns.kdeplot(df1['Flight Distance']);
```



```
[ ]: # check distribution of Departure Delay
plt.figure(figsize = (5,5))
sns.kdeplot(df1['Departure Delay']);
```



```
[ ]: # check distribution of Arrival Delay
plt.figure(figsize = (5,5))
sns.kdeplot(df1['Arrival Delay']);
```



Observation:

- Outliers in Flight Distance, Departure Delay and Arrival Delay
- The distribution of all three variables are right skewed (Positively Skewed). This means that the mean is often greater than the median.
- Arrival Delay includes 393 missing values.

## 4 Data preparation for Supervised Machine Learning

Split the original dataset (df1) into Training set and Test set

```
[88]: from sklearn.model_selection import train_test_split
```

```
[89]: X = df1.drop(['Satisfaction'], axis=1)
      y = df1['Satisfaction']
```

```
[90]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state= 0 ,
      ↪test_size=0.25, shuffle=True)
```

```
[91]: # view first few rows of train set
      X_train.head()
```

[91]: Gender Age Customer Type Type of Travel Class Flight Distance \

ID	Gender	Age	Customer Type	Type of Travel	Class	Flight Distance
2050	Male	47	Returning	Business	Economy	812
49177	Female	44	Returning	Business	Business	3285
38347	Male	26	Returning	Business	Economy	1173
36700	Male	48	Returning	Personal	Economy	1197
20522	Female	16	Returning	Personal	Business	533

ID	Departure Delay	Arrival Delay	Departure and Arrival Time Convenience
2050	63	51.0	3
49177	0	0.0	0
38347	0	0.0	3
36700	9	0.0	5
20522	0	0.0	5

ID	Ease of Online Booking	Gate Location	On-board Service
2050	3	3	1
49177	0	1	3
38347	3	3	3
36700	1	2	5
20522	1	4	4

ID	Seat Comfort	Leg Room Service	Cleanliness	Food and Drink
2050	3	1	3	3
49177	3	3	1	3
38347	4	5	5	5
36700	4	2	4	4
20522	5	5	5	5

ID	In-flight Service	In-flight Wifi Service	In-flight Entertainment
2050	3	3	3
49177	3	0	3
38347	4	5	5
36700	4	1	4
20522	5	1	5

ID	Baggage Handling
2050	3
49177	3
38347	4
36700	1
20522	5

[5 rows x 22 columns]

```
[92]: X_train["Arrival Delay"].isnull().sum()
```

[92]: 290

```
[93]: X_train["Arrival Delay"].mean()
```

[93]: 15.105889621087314

```
[94]: # view first few rows of test set
y_train.head()
```

[94]: ID  
2050        Neutral or Dissatisfied  
49177                    Satisfied  
38347                    Satisfied  
36700        Neutral or Dissatisfied  
20522        Neutral or Dissatisfied  
Name: Satisfaction, dtype: object

```
[95]: # check the size of each set
print(X_train.shape, X_test.shape)
```

(97410, 22) (32470, 22)

```
[96]: # check missing value of train set
X_train['Arrival Delay'].isnull().sum()
```

[96]: 290

```
[97]: # check missing value of test set
X_test['Arrival Delay'].isnull().sum()
```

[97]: 103


This test set contains some missing values. I am looking for a way to split the dataset into training and golden standard test set (no missing values in test set)

Handling missing value of train set

```
[98]: # check missing value of train set
missing = X_train.isnull().sum()
missing = missing[missing > 0]
missing = missing.sort_values(ascending = False)
missing
```

```
[98]: Arrival Delay    290  
      dtype: int64
```

```
[99]: # check the distribution of Arrival Delay of train set  
      plt.figure(figsize = (5,5))  
      sns.kdeplot(X_train['Arrival Delay']);
```



#### Observation:

- Arrival Delay is right skewed distribution. Therefore, we will impute the median to the missing values.


```
[100]: modified_X_train = X_train  
      modified_X_train['Arrival Delay'].fillna(modified_X_train['Arrival Delay'].  
      ↪median(), inplace=True)
```

```
[101]: modified_X_train['Arrival Delay'].isnull().sum()
```

```
[101]: 0
```



```
[102]: # check the distribution of Arrival Delay of modified_train set (AFTER inputing
        ↳ median value)
plt.figure(figsize = (5,5))
sns.kdeplot(modified_X_train['Arrival Delay']);
```



**Observation:** There is no significant change AFTER inputing the median value to fill the missing values.

### Handling missing value of test set

To avoid data leakage, median value from Train train set is imputed to test set.

```
[103]: modified_X_test = X_test
```

```
[104]: modified_X_test['Arrival Delay'].fillna(modified_X_train['Arrival Delay'].
        ↳ median(), inplace=True)
modified_X_test['Arrival Delay'].isnull().sum()
```

```
[104]: 0
```

```
[105]: modified_X_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 32470 entries, 125670 to 34201
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                32470 non-null  object
1   Age                                   32470 non-null  int64
2   Customer Type                         32470 non-null  object
3   Type of Travel                        32470 non-null  object
4   Class                                 32470 non-null  object
5   Flight Distance                       32470 non-null  int64
6   Departure Delay                       32470 non-null  int64
7   Arrival Delay                         32470 non-null  float64
8   Departure and Arrival Time Convenience 32470 non-null  int64
9   Ease of Online Booking                32470 non-null  int64
10  Check-in Service                      32470 non-null  int64
11  Online Boarding                       32470 non-null  int64
12  Gate Location                         32470 non-null  int64
13  On-board Service                      32470 non-null  int64
14  Seat Comfort                          32470 non-null  int64
15  Leg Room Service                      32470 non-null  int64
16  Cleanliness                           32470 non-null  int64
17  Food and Drink                        32470 non-null  int64
18  In-flight Service                     32470 non-null  int64
19  In-flight Wifi Service                32470 non-null  int64
20  In-flight Entertainment                32470 non-null  int64
21  Baggage Handling                      32470 non-null  int64
dtypes: float64(1), int64(17), object(4)
memory usage: 5.7+ MB

```

## Handling outliers of train set

Use IQR (Inter Quartile Range) to finding the outliers and cap the outliers

- capping: to replace the outlier values with a maximum or minimum capped value

## Arrival Delay

```


[106]: # before capping outliers
sns.boxplot( y="Arrival Delay", data = modified_X_train)

```

```

[106]: <Axes: ylabel='Arrival Delay'>

```



```
[107]: # IQR
Q1 = np.percentile(modified_X_train['Arrival Delay'], 25, method='midpoint')
Q3 = np.percentile(modified_X_train['Arrival Delay'], 75, method='midpoint')
IQR = Q3 - Q1
print(IQR)
```

13.0

```
[108]: upper_bound = Q3 + 1.5 * IQR
lower_bound = Q1 - 1.5 * IQR
print(upper_bound)
print(lower_bound)
```


32.5

-19.5

```
[109]: modified_X_train["Arrival Delay"] = np.where(modified_X_train["Arrival Delay"] > upper_bound, upper_bound,
                                                    np.where(modified_X_train["Arrival Delay"] < lower_bound, lower_bound,
                                                                modified_X_train["Arrival Delay"]))
```

```
[110]: # after capping outliers
sns.boxplot( y="Arrival Delay", data = modified_X_train)
```


```
[110]: <Axes: ylabel='Arrival Delay'>
```



### Flight Distance

```
[111]: # before capping outliers
sns.boxplot( y="Flight Distance", data = modified_X_train)
# sns.scatterplot(x= modified_train['Flight Distance'] , y =
↳ modified_train['Satisfaction'])
```

```
[111]: <Axes: ylabel='Flight Distance'>
```



```
[112]: # IQR
Q1 = np.percentile(df1['Flight Distance'], 25, method='midpoint')
Q3 = np.percentile(df1['Flight Distance'], 75, method='midpoint')
IQR = Q3 - Q1
print(IQR)
```

1330.0

```
[113]: upper_bound = Q3 + 1.5 * IQR
lower_bound = Q1 - 1.5 * IQR
print(upper_bound)
print(lower_bound)
```


3739.0

-1581.0

```
[114]: modified_X_train["Flight Distance"] = np.where(modified_X_train["Flight_
Distance"] > upper_bound, upper_bound,
               np.where(modified_X_train["Flight Distance"] <_
lower_bound, lower_bound,
               modified_X_train["Flight Distance"])))
```

```
[115]: # after capping outliers
sns.boxplot( y="Flight Distance", data = modified_X_train)
```


```
[115]: <Axes: ylabel='Flight Distance'>
```



### Departure Delay

```
[116]: # before capping outliers
sns.boxplot( y="Departure Delay", data = modified_X_train)
# sns.scatterplot(x= modified_train['Departure Delay'] , y =
↳ modified_train['Satisfaction'])
```

```
[116]: <Axes: ylabel='Departure Delay'>
```



```
[117]: # IQR
Q1 = np.percentile(df1['Departure Delay'], 25, method='midpoint')
Q3 = np.percentile(df1['Departure Delay'], 75, method='midpoint')
IQR = Q3 - Q1
print(IQR)
```

12.0

```
[118]: upper_bound = Q3 + 1.5 * IQR
lower_bound = Q1 - 1.5 * IQR
print(upper_bound)
print(lower_bound)
```


30.0

-18.0

```
[119]: modified_X_train["Departure Delay"] = np.where(modified_X_train["Departure_
↳ Delay"] > upper_bound, upper_bound,
               np.where(modified_X_train["Departure Delay"] <_
↳ lower_bound, lower_bound,
               modified_X_train["Departure Delay"]))
```

```
[120]: # after capping outliers
sns.boxplot( y="Departure Delay", data = modified_X_train)
```

```
[120]: <Axes: ylabel='Departure Delay'>
```



Encoding all categorical variables (columns with type : object) to numerical variables

```
[121]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
[122]: encoded_X_train = modified_X_train
```

```
[123]: encoded_X_train["Gender"].value_counts()
```

```
[123]: Female    49432
Male        47978
Name: Gender, dtype: int64
```

```
[124]: encoded_X_train["Gender"] = le.fit_transform(encoded_X_train["Gender"])
encoded_X_train["Gender"].value_counts()
```

```
[124]: 0    49432
1    47978
```



Name: Gender, dtype: int64

```
[125]: encoded_X_train["Customer Type"].value_counts()
```

```
[125]: Returning      79497
First-time      17913
Name: Customer Type, dtype: int64
```

```
[126]: encoded_X_train["Customer Type"]=le.fit_transform(encoded_X_train["Customer_
↪Type"])
encoded_X_train["Customer Type"].value_counts()
```

```
[126]: 1      79497
0      17913
Name: Customer Type, dtype: int64
```

```
[127]: encoded_X_train["Type of Travel"].value_counts()
```

```
[127]: Business      67234
Personal      30176
Name: Type of Travel, dtype: int64
```

```
[128]: encoded_X_train["Type of Travel"]=le.fit_transform(encoded_X_train["Type of_
↪Travel"])
encoded_X_train["Type of Travel"].value_counts()
```

```
[128]: 0      67234
1      30176
Name: Type of Travel, dtype: int64
```

```
[129]: encoded_X_train["Class"].value_counts()
```

```
[129]: Business      46509
Economy      43832
Economy Plus      7069
Name: Class, dtype: int64
```

```
[130]: encoded_X_train["Class"]=le.fit_transform(encoded_X_train["Class"])
encoded_X_train["Class"].value_counts()
```

```
[130]: 0      46509
1      43832
2       7069
Name: Class, dtype: int64
```

```
[131]: encoded_X_train.head()
```

[131]:

	Gender	Age	Customer Type	Type of Travel	Class	Flight Distance	\
ID							
2050	1	47	1	0	1	812.0	
49177	0	44	1	0	0	3285.0	
38347	1	26	1	0	1	1173.0	
36700	1	48	1	1	1	1197.0	
20522	0	16	1	1	0	533.0	

	Departure Delay	Arrival Delay	Departure and Arrival Time Convenience	\
ID				
2050	30.0	32.5		3
49177	0.0	0.0		0
38347	0.0	0.0		3
36700	9.0	0.0		5
20522	0.0	0.0		5

	Ease of Online Booking	...	Gate Location	On-board Service	\
ID		...			
2050	3	...	3	1	
49177	0	...	1	3	
38347	3	...	3	3	
36700	1	...	2	5	
20522	1	...	4	4	

	Seat Comfort	Leg Room Service	Cleanliness	Food and Drink	\
ID					
2050	3	1	3	3	
49177	3	3	1	3	
38347	4	5	5	5	
36700	4	2	4	4	
20522	5	5	5	5	

	In-flight Service	In-flight Wifi Service	In-flight Entertainment	\
ID				
2050	3	3	3	
49177	3	0	3	
38347	4	5	5	
36700	4	1	4	
20522	5	1	5	

	Baggage Handling
ID	
2050	3
49177	3
38347	4
36700	1
20522	5

[5 rows x 22 columns]

```
[132]: y_train.value_counts()
```

```
[132]: Neutral or Dissatisfied    55153  
      Satisfied                42257  
      Name: Satisfaction, dtype: int64
```

```
[133]: encoded_y_train=le.fit_transform(y_train)  
      encoded_y_train
```

```
[133]: array([0, 1, 1, ..., 1, 0, 0])
```

```
[134]: encoded_y_train = pd.DataFrame(encoded_y_train, index=encoded_y_train)
```

```
[135]: encoded_y_train.columns = ['Satisfaction']
```

```
[136]: encoded_y_train.head()
```

```
[136]:      Satisfaction  
      0            0  
      1            1  
      1            1  
      0            0  
      0            0
```

### Applying Feature Selection to reduce dimensions

```
[137]: # Pre-processing and scaling dataset for feature selection  
      from sklearn import preprocessing  
  
      r_scaler = preprocessing.MinMaxScaler()  
      r_scaler.fit(encoded_X_train)  
  
      encoded_X_train_scaled = pd.DataFrame(r_scaler.transform(encoded_X_train),  
      ↪ columns = encoded_X_train.columns)  
      encoded_X_train_scaled.head()  
  
      encoded_y_train_scaled = encoded_y_train
```

```
[138]: # Finding the best K for feature selection  
  
      import sklearn.feature_selection as fs  
      import sklearn.datasets as datasets  
      from sklearn.model_selection import train_test_split  
      from sklearn.linear_model import LogisticRegression
```

```

import sklearn.metrics as metrics
import matplotlib.pyplot as plt

X = encoded_X_train_scaled
y = encoded_y_train_scaled

f1_list = []
for k in range(1, 22):
    bk = fs.SelectKBest(fs.f_classif, k = k)
    bk.fit(X, y)
    X_trans = bk.transform(X)
    train_x, test_x, train_y, test_y = train_test_split(X_trans,
                                                         y,
                                                         test_size=0.2,
                                                         random_state=42)

    lr = LogisticRegression()
    lr.fit(train_x, train_y)
    y_pred = lr.predict(test_x)
    #f1 = metrics.f1_score(test_y, y_pred, pos_label="Satisfied")
    f1 = metrics.f1_score(test_y, y_pred)
    f1_list.append(f1)

print(len(f1_list))

fig, axe = plt.subplots(dpi = 150)

print(type(axe))

axe.plot(range(0, len(f1_list)), f1_list)
axe.set_xlabel("best k features")
axe.set_ylabel("F1-score")
plt.grid(True)
plt.show()
# fig.savefig("img.png")
# plt.close(fig)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().

```

```

    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().

```

```

    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:

```

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```

y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was

```

expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)


```



```

ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
21
<class 'matplotlib.axes._axes.Axes'>

```



```

[139]: # Feature selection, applying Select K Best to output the 14 most important_
        ↪ features
from sklearn.feature_selection import SelectKBest, f_classif

# X = encoded_X_train_scaled.loc[:, encoded_X_train_scaled.columns!
        ↪ 'Satisfaction']
X = encoded_X_train_scaled
y = encoded_y_train_scaled

selector = SelectKBest(f_classif, k = 14)
selector.fit(X, y)

```

```
selected_X_train_scaled = selector.transform(X)

features = (X.columns[selector.get_support(indices=True)])
features
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
```

```
[139]: Index(['Type of Travel', 'Class', 'Flight Distance', 'Check-in Service',
            'Online Boarding', 'On-board Service', 'Seat Comfort',
            'Leg Room Service', 'Cleanliness', 'Food and Drink',
            'In-flight Service', 'In-flight Wifi Service',
            'In-flight Entertainment', 'Baggage Handling'],
            dtype='object')
```

```
[140]: features = pd.DataFrame(selected_X_train_scaled[selector.
    ↪get_support(indices=True)])
    # features.head()
```

```
[141]: features.columns = ['Type of Travel', 'Class', 'Flight Distance', 'Check-in_
    ↪Service',
            'Online Boarding', 'On-board Service', 'Seat Comfort',
            'Leg Room Service', 'Cleanliness', 'Food and Drink',
            'In-flight Service', 'In-flight Wifi Service',
            'In-flight Entertainment', 'Baggage Handling']
```

```
[142]: features.head()
```

```
[142]:
```

	Type of Travel	Class	Flight Distance	Check-in Service	Online Boarding	\
0	1.0	0.5	0.314455	0.4	0.2	
1	1.0	0.0	0.135383	0.6	0.2	
2	0.0	0.5	0.174757	0.8	0.6	
3	1.0	0.5	0.154531	1.0	0.2	
4	0.0	0.0	0.307983	0.8	0.8	

	On-board Service	Seat Comfort	Leg Room Service	Cleanliness	\
0	1.0	0.8	0.4	0.8	
1	0.8	1.0	1.0	1.0	
2	0.8	0.4	0.2	0.4	
3	1.0	1.0	1.0	1.0	
4	1.0	1.0	0.8	1.0	

	Food and Drink	In-flight Service	In-flight Wifi Service	\
0	0.8	0.8	0.2	

1	1.0	1.0	0.2
2	0.4	0.6	0.6
3	1.0	1.0	0.2
4	1.0	0.8	0.8

	In-flight Entertainment	Baggage Handling
0	0.8	0.00
1	1.0	1.00
2	0.4	0.75
3	1.0	0.75
4	1.0	1.00

```
[143]: selected_X_train_scaled.shape
```

```
[143]: (97410, 14)
```

Applying undersampling, oversampling and SMOTE to address the imbalance in the target class, “Satisfaction”.

### Applying SMOTE

```
[144]: from imblearn.over_sampling import SMOTE
from collections import Counter
# define dataset
X_SMOTE = selected_X_train_scaled
y_SMOTE = encoded_y_train_scaled
# summarize class distribution
counter = Counter(y_SMOTE)
print('Before SMOTE', (counter))
# transform the dataset
oversample = SMOTE()
X_SMOTE, y_SMOTE = oversample.fit_resample(X_SMOTE, y_SMOTE)
# summarize the new class distribution
counter = Counter(y_SMOTE)
print('After SMOTE', (counter))
```

Before SMOTE Counter({'Satisfaction': 1})

After SMOTE Counter({'Satisfaction': 1})

```
[145]: SMOTE_train = pd.DataFrame(X_SMOTE)
SMOTE_train.columns = features.columns
```

```
[146]: SMOTE_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110306 entries, 0 to 110305
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
#   ...
```

```

---  -----
0   Type of Travel      110306 non-null float64
1   Class               110306 non-null float64
2   Flight Distance     110306 non-null float64
3   Check-in Service    110306 non-null float64
4   Online Boarding     110306 non-null float64
5   On-board Service    110306 non-null float64
6   Seat Comfort        110306 non-null float64
7   Leg Room Service    110306 non-null float64
8   Cleanliness         110306 non-null float64
9   Food and Drink      110306 non-null float64
10  In-flight Service   110306 non-null float64
11  In-flight Wifi Service 110306 non-null float64
12  In-flight Entertainment 110306 non-null float64
13  Baggage Handling     110306 non-null float64
dtypes: float64(14)
memory usage: 11.8 MB

```

```
[147]: y_SMOTE.value_counts()
```

```

[147]: Satisfaction
0          55153
1          55153
dtype: int64

```

## Applying undersampling

```

[148]: from imblearn.under_sampling import RandomUnderSampler
# define dataset
X_under = selected_X_train_scaled
y_under = encoded_y_train_scaled
# summarize class distribution
print('Before UnderSampling', (Counter(y_under)))
# define undersample strategy
undersample = RandomUnderSampler(sampling_strategy='majority')
# fit and apply the transform
X_under, y_under = undersample.fit_resample(X_under, y_under)
# summarize class distribution
print('After UnderSampling', (Counter(y_under)))

```

Before UnderSampling Counter({'Satisfaction': 1})

After UnderSampling Counter({'Satisfaction': 1})

```

[149]: undersampling_train = pd.DataFrame(X_under)
undersampling_train.columns = features.columns

```

```
[150]: undersampling_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 84514 entries, 0 to 84513
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Type of Travel                        84514 non-null  float64
1   Class                                84514 non-null  float64
2   Flight Distance                       84514 non-null  float64
3   Check-in Service                     84514 non-null  float64
4   Online Boarding                      84514 non-null  float64
5   On-board Service                     84514 non-null  float64
6   Seat Comfort                         84514 non-null  float64
7   Leg Room Service                     84514 non-null  float64
8   Cleanliness                          84514 non-null  float64
9   Food and Drink                       84514 non-null  float64
10  In-flight Service                    84514 non-null  float64
11  In-flight Wifi Service                84514 non-null  float64
12  In-flight Entertainment               84514 non-null  float64
13  Baggage Handling                     84514 non-null  float64
dtypes: float64(14)
memory usage: 9.0 MB
```

```
[151]: y_under.value_counts()
```

```
[151]: Satisfaction
0          42257
1          42257
dtype: int64
```

## Applying oversampling

```
[152]: from imblearn.over_sampling import RandomOverSampler
# define dataset
X_over = selected_X_train_scaled
# y_over = y_train
y_over = encoded_y_train_scaled
# summarize class distribution
print('Before OverSampling', (Counter(y_over)))
# define oversampling strategy
oversample = RandomOverSampler(sampling_strategy='minority')
# fit and apply the transform
X_over, y_over = oversample.fit_resample(X_over, y_over)
# summarize class distribution
print('After OverSampling', (Counter(y_over)))
```

```
Before OverSampling Counter({'Satisfaction': 1})
After OverSampling Counter({'Satisfaction': 1})
```

```
[153]: oversampling_train = pd.DataFrame(X_over)
oversampling_train.columns = features.columns
```

```
[154]: oversampling_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110306 entries, 0 to 110305
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Type of Travel                        110306 non-null float64
 1   Class                                110306 non-null float64
 2   Flight Distance                       110306 non-null float64
 3   Check-in Service                     110306 non-null float64
 4   Online Boarding                      110306 non-null float64
 5   On-board Service                     110306 non-null float64
 6   Seat Comfort                         110306 non-null float64
 7   Leg Room Service                     110306 non-null float64
 8   Cleanliness                          110306 non-null float64
 9   Food and Drink                       110306 non-null float64
10   In-flight Service                    110306 non-null float64
11   In-flight Wifi Service               110306 non-null float64
12   In-flight Entertainment              110306 non-null float64
13   Baggage Handling                     110306 non-null float64
dtypes: float64(14)
memory usage: 11.8 MB
```

```
[156]: y_over.value_counts()
```

```
[156]: Satisfaction
0          55153
1          55153
dtype: int64
```

## Building Models

Apply Random Forest, k-Nearest Neighbours, and Gradient Boosting (Extreme Gradient Boosting (XGBoost))

### Random Forest: SMOTE

```
[157]: from sklearn.ensemble import RandomForestClassifier
```

```
[158]: RF = RandomForestClassifier(max_features= 14, max_depth=7)
```

```
[159]: RF.fit(SMOTE_train , y_SMOTE)
RF.score(SMOTE_train , y_SMOTE)
```

```
<ipython-input-159-d365d61def11>:1: DataConversionWarning: A column-vector y was
```

passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
RF.fit(SMOTE_train , y_SMOTE)
```

[159]: 0.9291969611807155

### Random Forest: undersampling

```
[160]: RF.fit(undersampling_train, y_under)
       RF.score(undersampling_train, y_under)
```

<ipython-input-160-8503aaaec792>:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
RF.fit(undersampling_train, y_under)
```

[160]: 0.9274203090612206

### Random Forest: oversampling

```
[161]: RF.fit(oversampling_train, y_over)
       RF.score(oversampling_train, y_over)
```

<ipython-input-161-c7bc34fba153>:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
RF.fit(oversampling_train, y_over)
```

[161]: 0.9266495022936195

**Observation:** Random Forest: SMOTE is the highest score, 92.9 but there is no significant difference.

### Checking the important features using Random Forest: SMOTE

```
[180]: RF_SMOTE_train = SMOTE_train
       RF_y_SMOTE = y_SMOTE

def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center', color = 'LightBlue')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances for Random Forest: SMOTE')
    plt.show()
```

```


features_names = RF_SMOTE_train.columns

rf = RandomForestClassifier(n_estimators=4 , max_depth=3 , min_samples_split=25,
↪, max_features=4, random_state=0)
rf.fit(RF_SMOTE_train , RF_y_SMOTE)
f_importances(abs(rf.feature_importances_), features_names, top=7)

```

<ipython-input-180-7e9ccdfb3c27>:19: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
rf.fit(RF_SMOTE_train , RF_y_SMOTE)
```



## KNN: SMOTE

```
[169]: from sklearn.neighbors import KNeighborsClassifier
```

```
[170]: KNN = KNeighborsClassifier(n_neighbors=5)
```

```
[171]: KNN.fit(SMOTE_train , y_SMOTE)
KNN.score(SMOTE_train , y_SMOTE)
```

/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/\_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape



```
of y to (n_samples,), for example using ravel().
return self._fit(X, y)
```

[171]: 0.9552698855909924

### KNN: undersampling

```
[172]: KNN.fit(undersampling_train, y_under)
KNN.score(undersampling_train, y_under)
```

```
/usr/local/lib/python3.10/dist-
packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape
of y to (n_samples,), for example using ravel().
return self._fit(X, y)
```

[172]: 0.9443287502662281

### KNN: oversampling

```
[86]: KNN.fit(oversampling_train, y_over)
KNN.score(oversampling_train, y_over)
```

```
/usr/local/lib/python3.10/dist-
packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape
of y to (n_samples,), for example using ravel().
return self._fit(X, y)
```

[86]: 0.9511268652657153

**Observation:** KNN: SMOTE is the highest score, 95.5 but there is no significant difference.

### Checking the important features using KNN: SMOTE

```
[209]: KNN_SMOTE_train = SMOTE_train
KNN_y_SMOTE = y_SMOTE

def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center', color = 'LightBlue')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances for KNN: SMOTE ')
    plt.show()
```

```

features_names = KNN_SMOTE_train.columns

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(KNN_SMOTE_train , KNN_y_SMOTE)
f_importances(abs(rf.feature_importances_), features_names, top=8)


```

/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/\_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```

return self._fit(X, y)

```



### XGBoost: SMOTE

```
[174]: from xgboost import XGBClassifier
```

```
[175]: XGB = XGBClassifier(max_depth = 18 , n_estimators= 6)
```

```
[176]: XGB.fit(SMOTE_train , y_SMOTE)
XGB.score(SMOTE_train , y_SMOTE)
```

```
[176]: 0.9748608416586586
```

### XGBoost: undersampling

```
[177]: XGB.fit(undersampling_train, y_under)
XGB.score(undersampling_train, y_under)
```

```
[177]: 0.9741107982109473
```

### XGBoost: oversampling

```
[178]: XGB.fit(oversampling_train, y_over)
XGB.score(oversampling_train, y_over)
```

```
[178]: 0.9762388265370877
```

**Observation:** XGBoost: oversampling = Highest score, 97.6 but but there is no significant difference.

### Checking the important features using XGBoost: oversampling

```
[210]: XGB_oversampling_train = oversampling_train
XGB_y_over = y_over


def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center', color = 'LightBlue')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances for XGBoost: oversampling ')
    plt.show()

features_names = XGB_oversampling_train.columns

xgb = XGBClassifier(max_depth = 18 , n_estimators= 6)
xgb.fit(XGB_oversampling_train , XGB_y_over)
f_importances(abs(xgb.feature_importances_), features_names, top=8)
```



```
[213]: import sklearn
import time
from resource import getrusage, RUSAGE_SELF
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
```

```
[216]: # Model activation and result plot function
def get_model_metrics(model, SMOTE_train, modified_X_test, y_SMOTE, y_test):

    """
    Model activation function, takes in model as a parameter and returns
    metrics as specified.

    Inputs:
        model, SMOTE_train, modified_X_test, y_SMOTE, y_test
    Output:
        Model output metrics, confusion matrix, ROC AUC curve
    """

    # Mark of current time when model began running
```

```

t0 = time.time()

# Fit the model on the training data and run predictions on test data
model.fit(SMOTE_train, y_SMOTE)
y_pred = model.predict(modified_X_test)
y_pred_proba = model.predict_proba(modified_X_test)[:,-1]
# Obtain training accuracy as a comparative metric using Sklearn's metrics
↪package
train_score = model.score(SMOTE_train, y_SMOTE)
# Obtain testing accuracy as a comparative metric using Sklearn's metrics
↪package
accuracy = accuracy_score(y_test, y_pred)
# Obtain precision from predictions using Sklearn's metrics package
precision = precision_score(y_test, y_pred)
# Obtain recall from predictions using Sklearn's metrics package
recall = recall_score(y_test, y_pred)
# Obtain ROC score from predictions using Sklearn's metrics package
roc = roc_auc_score(y_test, y_pred_proba)
# Obtain the time taken used to run the model, by subtracting the start
↪time from the current time
time_taken = time.time() - t0
# Obtain the resources consumed in running the model
memory_used = int(gettrusage(RUSAGE_SELF).ru_maxrss / 1024)

# Outputting the metrics of the model performance
print("Accuracy on Training = {}".format(train_score))
print("Accuracy on Test = {} • Precision = {}".format(accuracy, precision))
print("Recall = {} • ROC Area under Curve = {}".format(recall, roc))
print("F1 = {} • ROC Area under Curve = {}".format(f1, roc))
print("Time taken = {} seconds • Memory consumed = {} Bytes".
↪format(time_taken, memory_used))

# Plotting the confusion matrix of the model's predictive capabilities
plt.confusion_matrix(model, modified_X_test, y_test, cmap = plt.cm.Blues,
↪normalize = 'all')
# Plotting the ROC AUC curve of the model
plt.roc_curve(model, modified_X_test, y_test)
plt.show()

return model, train_score, accuracy, precision, recall, roc, time_taken,
↪memory_used

```